

kaggle new york city taxi fare prediction challenge

September 25, 2018

```
In [1]: # Initial Python environment setup...
import numpy as np # linear algebra
import pandas as pd # CSV file I/O (e.g. pd.read_csv)
import os # reading the input files we have access to

print(os.listdir('../input'))

['GCP-Coupons-Instructions.rtf', 'sample_submission.csv', 'test.csv', 'train.csv']
```

```
In [2]: train_df = pd.read_csv('../input/train.csv',nrows=10_000_000)
train_df.dtypes
```

```
Out[2]: key                object
fare_amount              float64
pickup_datetime          object
pickup_longitude         float64
pickup_latitude          float64
dropoff_longitude         float64
dropoff_latitude          float64
passenger_count           int64
dtype: object
```

```
In [3]: #get a sense how training data looks like
train_df[:5]
```

```
Out[3]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21 UTC	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16 UTC	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00 UTC	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42 UTC	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00 UTC	

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	

4	-73.968095	40.768008	-73.956655	40.783762
---	------------	-----------	------------	-----------

	passenger_count
0	1
1	1
2	2
3	1
4	1

```
In [4]: plot = train_df.plot.scatter('passenger_count', 'fare_amount')
```

```
In [5]: train_df.describe()
```

```
Out [5]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	1.000000e+07	1.000000e+07	1.000000e+07	9.999931e+06	
mean	1.133854e+01	-7.250775e+01	3.991934e+01	-7.250897e+01	
std	9.799930e+00	1.299421e+01	9.322539e+00	1.287532e+01	
min	-1.077500e+02	-3.439245e+03	-3.492264e+03	-3.426601e+03	
25%	6.000000e+00	-7.399207e+01	4.073491e+01	-7.399139e+01	
50%	8.500000e+00	-7.398181e+01	4.075263e+01	-7.398016e+01	
75%	1.250000e+01	-7.396710e+01	4.076712e+01	-7.396367e+01	
max	1.273310e+03	3.457626e+03	3.344459e+03	3.457622e+03	

	dropoff_latitude	passenger_count
count	9.999931e+06	1.000000e+07
mean	3.991913e+01	1.684793e+00
std	9.237280e+00	1.323423e+00
min	-3.488080e+03	0.000000e+00
25%	4.073403e+01	1.000000e+00
50%	4.075316e+01	1.000000e+00
75%	4.076810e+01	2.000000e+00
max	3.351403e+03	2.080000e+02

```
In [6]: # Given a dataframe, add two new features 'abs_diff_longitude' and
# 'abs_diff_latitude' representing the "Manhattan vector" from
# the pickup location to the dropoff location.
def add_travel_vector_features(df):
    df['abs_diff_longitude'] = (df.dropoff_longitude - df.pickup_longitude).abs()
    df['abs_diff_latitude'] = (df.dropoff_latitude - df.pickup_latitude).abs()

add_travel_vector_features(train_df)
```

```
In [7]: #Clean NA values
print('Old size: %d' % len(train_df))
train_df = train_df.dropna(how = 'any', axis = 'rows')
print('New size: %d' % len(train_df))
```

```
Old size: 10000000
New size: 9999931
```

```
In [8]: print('Old size: %d' % len(train_df))
        train_df = train_df[(train_df.abs_diff_longitude < 5.0) & (train_df.abs_diff_latitude < 5.0)]
        print('New size: %d' % len(train_df))
```

Old size: 9999931

New size: 9979187

```
In [9]: print('Old size: %d' % len(train_df))
        train_df = train_df[(train_df.fare_amount > 0)]
        print('New size: %d' % len(train_df))
```

Old size: 9979187

New size: 9978549

```
In [10]: print('Old size: %d' % len(train_df))
          train_df = train_df[(train_df.pickup_longitude != 0) & (train_df.pickup_latitude != 0)]
          print('New size: %d' % len(train_df))
```

Old size: 9978549

New size: 9797235

```
In [11]: print('Old size: %d' % len(train_df))
          train_df = train_df[(train_df.dropoff_longitude != 0) & (train_df.dropoff_latitude != 0)]
          print('New size: %d' % len(train_df))
```

Old size: 9797235

New size: 9797221

```
In [12]: # Given a dataframe, add two new features 'abs_diff_longitude' and
          # 'abs_diff_latitude' representing the "Manhattan vector" from
          # the pickup location to the dropoff location.
          def add_euclid_distance_feature(df):
              abs_diff_longitude = (df.dropoff_longitude - df.pickup_longitude)
              abs_diff_latitude = (df.dropoff_latitude - df.pickup_latitude)
              abs_diff_longitude_square = abs_diff_longitude ** 2
              abs_diff_latitude_square = abs_diff_latitude ** 2
              df['euclid_dist'] = (abs_diff_longitude_square + abs_diff_latitude_square) ** (1/2)

          add_euclid_distance_feature(train_df)
```

```
In [13]: train_df.corr()
```

```
Out[13]:
```


	fare_amount	pickup_longitude	pickup_latitude	\
fare_amount	1.000000	0.005782	-0.004231	
pickup_longitude	0.005782	1.000000	-0.353907	

pickup_latitude	-0.004231	-0.353907	1.000000
dropoff_longitude	0.004879	0.999945	-0.353907
dropoff_latitude	-0.004082	-0.353901	0.999954
passenger_count	0.013833	0.005445	-0.005951
abs_diff_longitude	0.710765	0.002789	-0.000495
abs_diff_latitude	0.536288	0.003477	-0.003181
euclid_dist	0.726546	0.003378	-0.001911

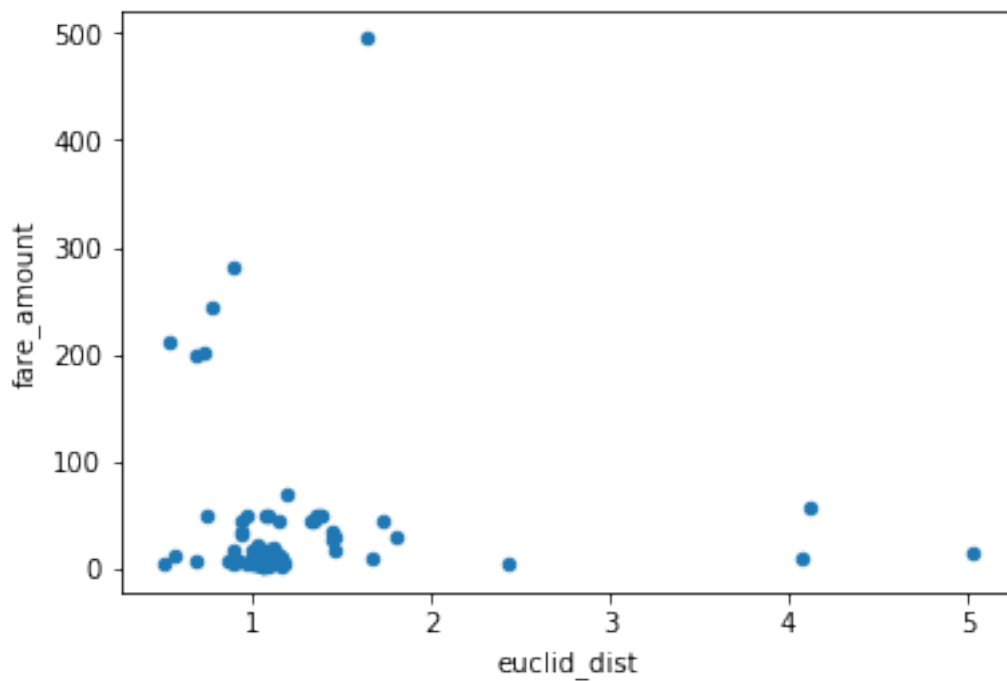
	dropoff_longitude	dropoff_latitude	passenger_count	\
fare_amount	0.004879	-0.004082	0.013833	
pickup_longitude	0.999945	-0.353901	0.005445	
pickup_latitude	-0.353907	0.999954	-0.005951	
dropoff_longitude	1.000000	-0.353897	0.005428	
dropoff_latitude	-0.353897	1.000000	-0.005928	
passenger_count	0.005428	-0.005928	1.000000	
abs_diff_longitude	0.002041	-0.000392	0.007301	
abs_diff_latitude	0.002850	-0.003186	0.006233	
euclid_dist	0.002674	-0.001905	0.007672	

	abs_diff_longitude	abs_diff_latitude	euclid_dist
fare_amount	0.710765	0.536288	0.726546
pickup_longitude	0.002789	0.003477	0.003378
pickup_latitude	-0.000495	-0.003181	-0.001911
dropoff_longitude	0.002041	0.002850	0.002674
dropoff_latitude	-0.000392	-0.003186	-0.001905
passenger_count	0.007301	0.006233	0.007672
abs_diff_longitude	1.000000	0.507669	0.904494
abs_diff_latitude	0.507669	1.000000	0.808833
euclid_dist	0.904494	0.808833	1.000000


```
In [14]: plot = train_df.iloc[:10000].plot.scatter('euclid_dist', 'fare_amount')
```



```
In [15]: #When euclid distance is more than 0.5, you can still fit a line but most of the data
x = train_df[train_df.euclid_dist>0.5]
plot = x.iloc[:100].plot.scatter('euclid_dist','fare_amount')
```



```
In [16]: #When euclid distance is less than 0.5, one can fit a line such that most data points
x = train_df[train_df.euclid_dist<0.5]
plot = x.iloc[:10000].plot.scatter('euclid_dist','fare_amount')
```



```
In [17]: #Convert pickup date time string to a pandas datetime object

from datetime import timedelta, datetime
from pandas import DataFrame, Series

lines = train_df['pickup_datetime']

dt_lst = []

for date_str in lines:
    dt = datetime.strptime(date_str, '%Y-%m-%d %H:%M:%S UTC')
    seconds_since_midnight = int((dt - dt.replace(hour=0, minute=0, second=0, microsecond=0)).total_seconds())
    dt_lst.append(seconds_since_midnight)

# Create a Series named "Request_Time"
sr_dt = Series(dt_lst, name='Second_Of_Day')

# Create a DataFrame using the Request_Time Series
```


```
df = DataFrame(sr_dt)
df['fare_amount'] = train_df['fare_amount']
df['euclid_dist'] = train_df['euclid_dist']
```

```
In [18]: df.corr()
```


```
Out[18]:
```

	Second_Of_Day	fare_amount	euclid_dist
Second_Of_Day	1.000000	-0.000365	-0.000165
fare_amount	-0.000365	1.000000	0.725529
euclid_dist	-0.000165	0.725529	1.000000

```
In [19]: plot = df.iloc[:10000].plot.scatter('Second_Of_Day', 'fare_amount')
```



```
In [20]: plot = df.iloc[:10000].plot.scatter('Second_Of_Day', 'euclid_dist')
```



```
In [21]: # Model 1 : A simple linear regression model that uses the feature euclidean distance
import matplotlib.pyplot as plt
from scipy import stats

xi = train_df['euclid_dist']
y = train_df['fare_amount']

#Train the model on train data
slope, intercept, r_value, p_value, std_err = stats.linregress(xi,y)

print ('r value', r_value)
print ('p_value', p_value)
print ('standard deviation', std_err)
print ('slope is:', slope)
print ("r-squared:", r_value**2)

#Check the model fit
#plt.plot(xi, y, 'o', label='original data')
#plt.plot(xi, intercept + slope*xi, 'r', label='fitted line')
#plt.legend()
#plt.show()

r value 0.7265459653520138
p_value 0.0
standard deviation 0.04594354813355792
```


slope is: 152.05748221076286
r-squared: 0.5278690397692896

```
In [22]: #Test the model on test data
test_df = pd.read_csv('../input/test.csv')
add_euclid_distance_feature(test_df)

xi = test_df['euclid_dist']

test_y_predictions=[]
for x in xi:
    y = round(slope*x + intercept, 2)
    test_y_predictions.append(y)

# Write the predictions to a CSV file which we can submit to the competition.
submission = pd.DataFrame(
    {'key': test_df.key, 'fare_amount': test_y_predictions},
    columns = ['key', 'fare_amount'])
submission.to_csv('submission.csv', index = False)

print(os.listdir('.'))
```

['.git', '.ipynb_checkpoints', 'kaggle new york city taxi fare prediction challenge.ipynb', 'mo

```
In [23]: # Model 2 : Use random forest regressor model that uses the feature euclidean distance
from sklearn.ensemble import RandomForestRegressor as rfr
from sklearn.model_selection import cross_val_score
# build the model with the desired parameters...
numFeatures = 1 # the number of features to include
train_df_mini = train_df[:1000000]
trees = 10 # trees in the forest
included_features = ['euclid_dist']
# define the training data X...
X = train_df_mini[included_features]
Y = train_df_mini[['fare_amount']]
yt = [i for i in Y['fare_amount']]
np.random.seed(11111)
model = rfr(n_estimators=trees,max_depth=None)
scores_rfr = cross_val_score(model,X,yt,cv=10,scoring='explained_variance')
print('explained variance scores for k=10 fold validation:',scores_rfr)
print("Est. explained variance: %0.2f (+/- %0.2f)" % (scores_rfr.mean(), scores_rfr.s
# fit the model
model.fit(X,yt)
```

explained variance scores for k=10 fold validation: [0.66106363 0.68731691 0.64530205 0.668022
0.70530754 0.69416429 0.69226933 0.6774113]

Est. explained variance: 0.68 (+/- 0.03)

```
Out [23]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
In [24]: # apply the model to the test data and get the output...
X_test = test_df[included_features]
y_output = model.predict(X_test.fillna(0)) # get the results and fill nan's with 0
print(y_output)

test_y_predictions=[]
for elem in y_output:
    test_y_predictions.append(round(elem,2))

[ 8.44973482  9.0342475   4.74166667 ... 42.648         21.472
 6.15284377]
```

```
In [25]: # Write the predictions to a CSV file which we can submit to the competition.
submission = pd.DataFrame(
    {'key': test_df.key, 'fare_amount': test_y_predictions},
    columns = ['key', 'fare_amount'])
submission.to_csv('submission.csv', index = False)

print(os.listdir('.'))
```

```
['.git', '.ipynb_checkpoints', 'kaggle new york city taxi fare prediction challenge.ipynb', 'm
```

```
In [26]: # Model 3 : Use LightGBM decision tree model that uses multiple features to predict t
import lightgbm as lgb
# load or create your dataset
print('Load data...')

included_features = ['euclid_dist', 'passenger_count', 'pickup_latitude', 'pickup_longitude']
train_X = train_df[included_features]
train_y = train_df['fare_amount']

params = {
    'nthread': -1,
}
#Default Parameters

train_set = lgb.Dataset(train_X, train_y, silent=True)
model3 = lgb.train(params, train_set = train_set, num_boost_round=300)
```

Load data...

```
In [27]: model3.save_model('model1lgbm.txt')
         print('Start predicting...')
         X_test = test_df[included_features]
         # predict
         y_pred = model3.predict(X_test, num_iteration=model3.best_iteration)
```

Start predicting...

```
In [28]: test_y_predictions=[]
         for elem in y_pred:
             test_y_predictions.append(round(elem,2))

         # Write the predictions to a CSV file which we can submit to the competition.
         submission = pd.DataFrame(
             {'key': test_df.key, 'fare_amount': test_y_predictions},
             columns = ['key', 'fare_amount'])
         submission.to_csv('submission.csv', index = False)

         print(os.listdir('.'))
```

```
['.git', '.ipynb_checkpoints', 'kaggle new york city taxi fare prediction challenge.ipynb', 'm
```